

# Process spawning and synchronisation basics.

Jiří Zahradník

Digiteq Automotive

2024

# What is a process

- General definition
  - A series of actions or steps taken in order to achieve a particular end.
- Definition in computing
  - An instance of a computer program, executed by one or more threads.

Both POSIX and QNX use so called POSIX thread model AKA pthreads. This contains API for:

- Thread management
- Mutexes
- Condition variables
- Synchronisation primitives (RW lock, barriers)
- Spinlocks (excellent for quick thread swapping, bad for prolonged waits)

# Some synchronisation methods and it's uses

- Spinlock
- Mutex
- Semaphore (Not included in pthread)
- Condition variable
- Monitor

# Spinlock - Simplest synchronisation primitive.

Uses active waiting loop and atomic operations upon an atomic flag. E.g. `test_and_set`, `clear` upon `std::atomic_flag`

## Pros

- Simple implementation.
- Lightweight implementation - does not take lot of instructions to perform.
- Good for very short critical sections and quick thread swapping.

# Spinlock - Simplest synchronisation primitive.

## Cons

- Uses busy loop wait.
- Bad for prolonged waits, since the thread is active, takes resources from other threads.
- Does not synchronise order of operations.
- To use in interprocess blocking, must be created in shared memory - not used unless necessary.

# Spinlock - example

```
class SpinLock {  
    std::atomic_flag locked = ATOMIC_FLAG_INIT ;  
public:  
    void lock() {  
        while (locked.test_and_set(  
            std::memory_order_acquire  
        )  
        ) { ; }  
    }  
    void unlock() {  
        locked.clear(std::memory_order_release);  
    }  
};
```

# Spinlock - example

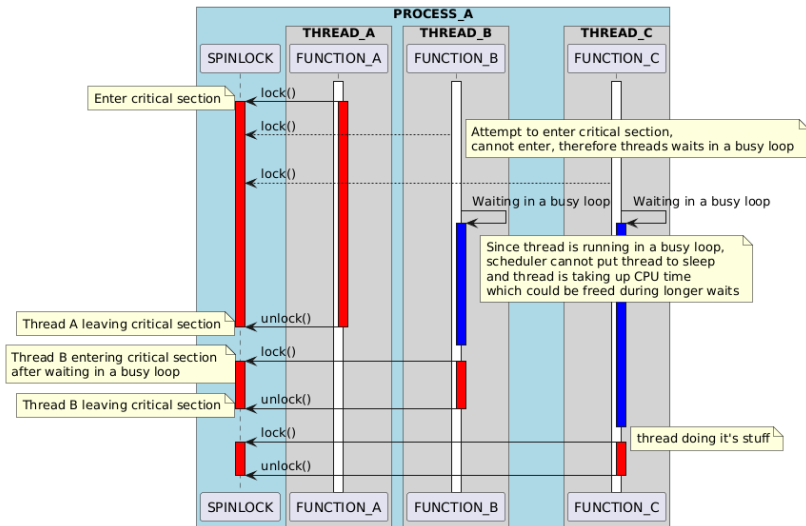


Figure: Spinlock sequence diagram



# Mutex

Mutual exclusion mechanism to secure critical section.

## Pros

- Easy usage with excellent abstraction.
- Easy usage in single process, multi threaded applications.
- Goes well with condition variable.
- (C++) Is accepted by various objects like `std::lock_guard`, `std::unique_lock`, `std::shared_lock`, `std::lock` etc.

## Cons

- Only owning thread can unlock the mutex, unlocking in different thread leads to undefined behaviour.
- A bit bigger overhead, but still nothing that should concern us unless we are building real time low power embedded systems.
- For interprocess communication requires special regime.

# Mutex - example

```
// entering critical section  
mutex.lock(); // if mutex is locked, thread blocks  
... // do stuff with shared resource  
mutex.unlock();
```

## Mutex - better example utilizing RAI1 in C++

```
foo function() {  
    // entering critical section  
    // Resource Allocation Is Initialisation.  
    // Acquiring mutex locks it  
    // and going out of scope unlocks it  
    // in destructor.  
    std::lock_guard<std::mutex> lock(mutex);  
  
    ... // do stuff with shared resource  
  
} // destructor releases the mutex (unlocks)
```

# Mutex - example

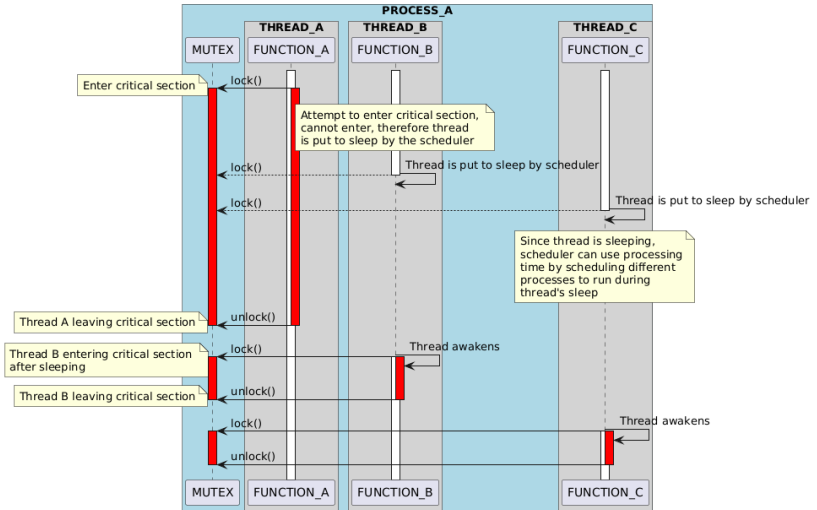


Figure: Mutex sequence diagram

# Mutex - `std::shared_mutex`

Contains 2 locks: `exclusive` and `shared`

- Acquiring exclusive lock prevents threads from acquiring both shared and exclusive locks.
- Acquiring shared lock prevents acquiring the exclusive lock, but allows acquiring more shared locks.
- Great solution for readers - writers problem.

# Semaphore

- Signaling primitive to synchronise multiple threads and processes.
- Consists of a counter, mutex and wait queue.
- Binary semaphore vs generic semaphore.
- Named semaphore vs unnamed semaphore.
- 2 operations
  - P() acquires semaphore (e.g. `sem_wait()`).
  - V() releases semaphore (e.g. `sem_post()`).
- A semaphore, by definition, does not have an owner.

## IMPORTANT!

Do not mistake binary semaphore for mutex.

## Pros

- Does not have an owner. Therefore one thread can wait, second can post.
- Easy to store named semaphore in shared memory to ensure interprocess communication.
- Named semaphore can synchronise two completely different processes which do not have to share same ancestor.
- We can decide between binary semaphore or general semaphore.
- Can protect a resource from being used by more than N processes/threads.
- Native C++ support as of C++20

## Cons

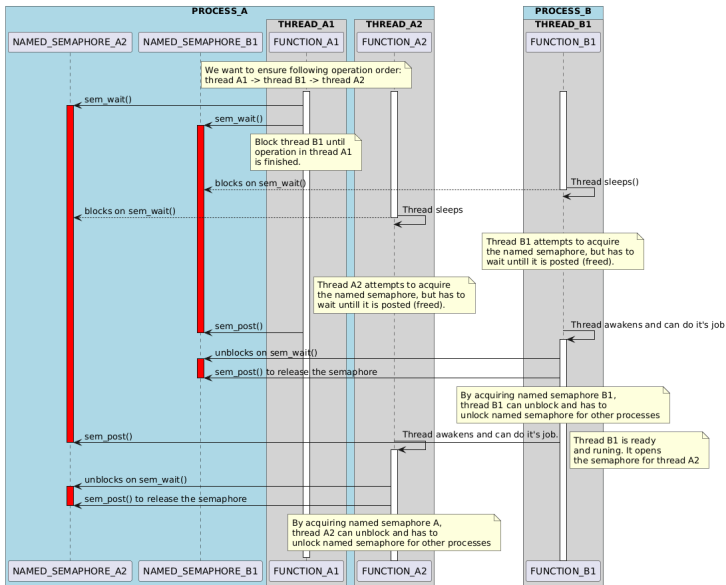
- Complexity - challenging to use correctly.
- Priority inversion - tasks with greater priority can be superseded by tasks with lower priority - important in real time tasks.
- They are expensive - CPU time and memory requirements.
- Requirement for correct access control list.

Priority inversion almost ended the Mars Pathfinder mission in 1977. More [here](#).

More on how to solve priority inversion later in the presentation.



# Semaphore example



# Condition variable

- Synchronisation primitive to enable threads to wait until a certain condition is satisfied.
- Used with mutexes to protect critical section.
- Primary use is in multithreading, not multiprocessing.
- Basically works as a queue for waiting threads.
- Is crucial for construction of monitors.
- In some cases can be better than semaphore.

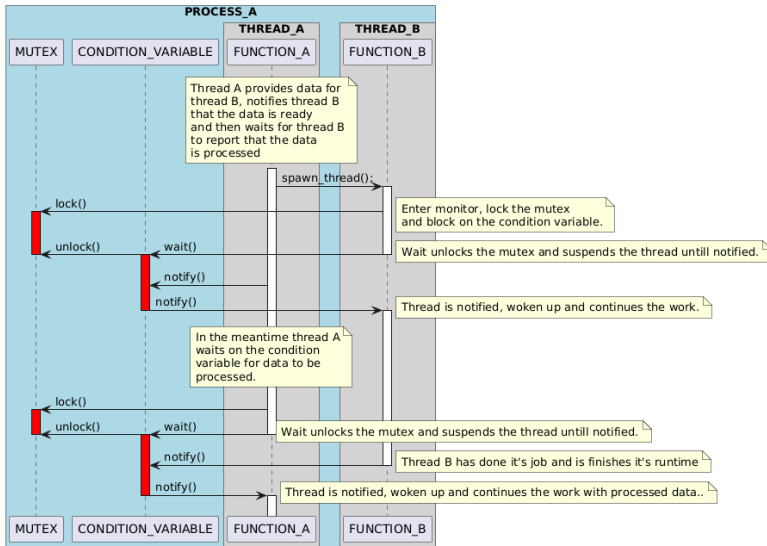
## Pros

- Simple synchronisation of threads.
- Combined with mutex solves consumer/producer problem.
- Native use in C/C++.
- User defined condition for wakeup.
- Can be used to wake up all threads.
- Solves spurious wakeups.

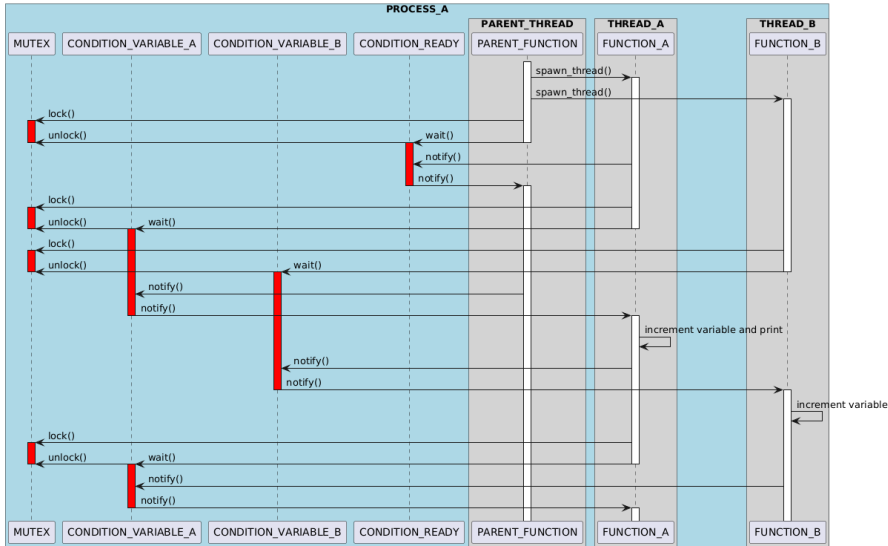
## Cons

- Can be only used in monitors.
- Primarily for multithreading synchronisation, multiprocess synchronisation requires further configuration and must be paired with shared mutex.
- `std::condition_variable` available only for multithreading.
- C version also available for multiprocessing, but requires further configuration.


# Condition variable example



# Condition variable vs Semaphore for synchronisation



# Condition variable vs Semaphore for synchronisation



- High level abstraction.
- Protection and synchronisation mechanism.
- Great for encapsulating protected data.
- Uses Mutex and at least one Condition variable to protect data inside.

## KEEP IN MIND

Rules still apply for mutex and condition variable.



# Spawning a new process.

## Common ways to spawn a process

- `fork()`
- `fork() + exec()`
- `clone()`
- `clone() + exec()`
- `posix_spawn() == clone() + exec()`

# Spawning a new process.

## Fork() and fork() + exec()

- `fork()` clones entire process using Copy on Write method, including parent process' memory map - in reality it's `clone()` with small preset overhead.
- `exec()` overwrites old memory map with new binary's memory map.

# Spawning a new process.

## `clone()` and `clone() + exec()`

- `clone()` permits more precise control over what pieces of execution context are shared.

# Spawning a new process.

## `spawn_posix()`

- wrapper for `clone()` and `exec()` with advantages and disadvantages.
  - Pros
    - Scheduling configuration passed as an argument
  - Cons
    - AFAIK core affinity configuration not supported (I didnt find it in documentation)

# Spawning a new thread

OS creates new execution context within a running process which shares code section and memory. Stack, registers including program counter are sole property of the thread and are not shared. This is called Thread Control Block

Possible issues.

# Deadlock

## Definition

Deadlock is any situation in which no member of some group of entities can proceed because each waits for another member, including itself, to take action, such as sending a message or, more commonly, releasing a lock.

## Typical representation

E.g. when two diggers share single shovel, but digger A holds the shovel, but can start digging after digger B finishes his hole, for which digger B needs the exact shovel.

## Invalidate Coffman conditions

For deadlock to occur, following necessary conditions must be satisfied.

- Mutual exclusion.
- No preemption.
- Hold and wait.
- Circular wait.



# Race condition

## Definition

Condition where the system's substantive behaviour is dependent on the sequence or timing of other uncontrollable events, leading to undefined behaviour.

## Typical representation

Data race during reader/writer problem.

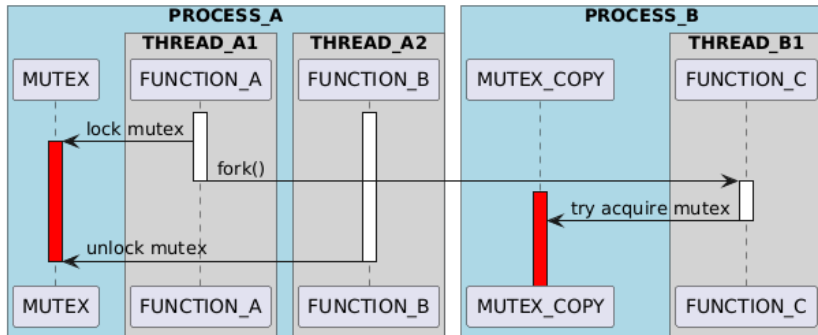
# Priority inversion

Priority inversion is a bug that occurs when a high priority task is indirectly preempted by a low priority task. For example, the low priority task holds a mutex that the high priority task must wait for to continue executing.

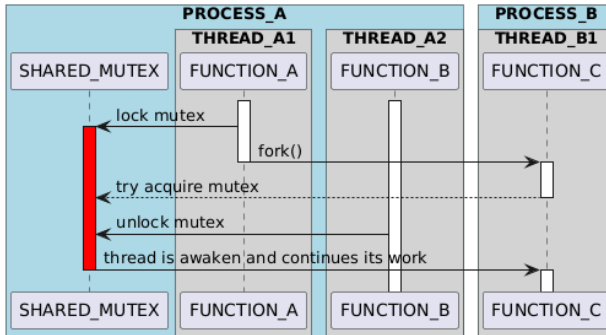
## Solutions

- Disabling all interrupts.
- Priority ceiling
- Priority Inheritance
- etc.

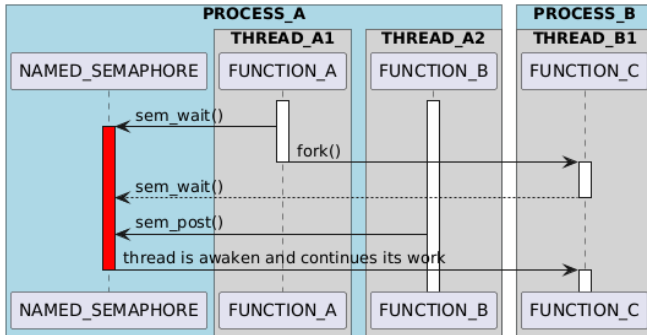
# Using wrong synchronisation primitive pt. 1



# Using wrong synchronisation primitive pt. 2



# Using correct synchronisation primitive pt. 1



# Using correct synchronisation primitive pt. 2

